

CPSC 223 Spring 2025 Exam #2 with Solutions

Tuesday, April 8, 2025

Name: _____ NetID: _____

Problem 1 (9 points): Hashtables

- a) Given an initially empty hashtable of capacity 4 using open addressing with linear probing and a resize threshold $\alpha = 0.5$, mark an X for the operation(s) that trigger a resize:

insert 10

insert 22

X insert 31

insert 4

insert 15

Explanation (a): The load factor is the number of entries divided by capacity. After inserting 10 and 22, $\text{load} = 2/4 = 0.5$, which is at threshold but does not exceed it. Inserting 31 makes $\text{load} = 3/4 = 0.75 > 0.5$, so the table is resized at that insertion.

- b) Suppose we delete key 22 and then search for key 15. Mark an X for each statement that is guaranteed true:

Search finds key 15 without probing past its initial slot.

X Search may require probing past deleted slots.

X Search fails if a tombstone is not handled.

Explanation (b): Deletion leaves a tombstone at 22's slot. If 15 hashed to a later slot, the search must probe past the tombstone to find it. Without special handling of tombstones, the probe might stop early and wrongly report 'not found.'

- c) What assumption about the hash function ensures average-case $O(1)$ performance?

Uniform hashing (each key equally likely to map to any slot, independently).

Explanation (c): Uniform hashing ensures that, on average, probes are equally likely to land anywhere, keeping the expected probe sequence short and thus guaranteeing constant-time performance in expectation.

Problem 2 (18 points): Binary Search Trees

For each code fragment below, write the number(s) that correctly build a BST containing keys [5, 2, 8, 1, 3] when *inserted in that order*.

1.

```
Node *root = NULL;
insert(&root, 5);
insert(&root, 2);
insert(&root, 8);
insert(&root, 1);
insert(&root, 3);
```

Explanation (2.1): This uses the standard insert function in the given order, so it builds the correct BST.

2.

```
Node *root = create_node(5);
root->left = create_node(2);
root->left->left = create_node(1);
root->left->right = create_node(3);
root->right = create_node(8);
```

Explanation (2.2): Nodes are manually linked in the exact BST structure matching keys 5 as root, 2–1–3 on the left, and 8 on the right.

3.

```
Node *arr[5];
arr[0]=create_node(5);
arr[1]=create_node(2);
arr[2]=create_node(8);
arr[3]=create_node(1);
arr[4]=create_node(3);
```

Explanation (2.3): Nodes are created but never linked; no tree is formed.

4.

```
Node *root = create_node(5);
insert(&root, 8);
insert(&root, 2);
insert(&root, 1);
insert(&root, 3);
```

Explanation (2.4): Inserts in the wrong order: 8 before 2 will place 8 as the right child first, then inserting 2 will still go left, but the structure/order of insert operations differs from the required sequence.

5.

```
Node *root = create_node(5);  
insert(&root, 2);  
insert(&root, 8);  
insert(&root, 3);  
insert(&root, 1);
```

Explanation (2.5): Swapping the order of 1 and 3 changes the placement: 3 is inserted before 1, altering the subtree shape incorrectly.

Answer: 1 and 2. **Explanation (2 overall):** Only fragments (1) and (2) produce the correct BST when keys [5, 2, 8, 1, 3] are inserted in that specified order. The others either never link properly or violate the insertion sequence.

Problem 3 (12 points): Graph Search and Shortest Paths

Consider the undirected graph G with adjacency lists:

$$A : B, C; \quad B : A, D, E; \quad C : A, F; \quad D : B; \quad E : B, F; \quad F : C, E.$$

- a) BFS from A (alphabetical neighbors): A, B, C, D, E, F

Explanation (3.a): BFS visits nodes level by level. From A, neighbors B and C come first; then from B we enqueue D and E; from C we enqueue F.

- b) DFS preorder from A: A, B, D, E, F, C

Explanation (3.b): Preorder visits the node, then recursively explores its neighbors in alphabetical order.

- c) DFS postorder from A: D, C, F, E, B, A

Explanation (3.c): Postorder explores children first, then visits the node after all deeper calls return.

- d) Dijkstra distances from A (unweighted, weight=1):

$$\text{dist}(A) = 0, \text{dist}(B) = 1, \text{dist}(C) = 1, \text{dist}(D) = 2, \text{dist}(E) = 2, \text{dist}(F) = 2$$

Explanation (3.d): All edges cost 1, so the distance is simply the minimum number of hops from A.

- e) Cycle? Yes

Explanation (3.e): A–B–E–F–C–A forms a cycle.

- f) Connected? Yes

Explanation (3.f): Every vertex is reachable from A.

Problem 4 (16 points): k-d Tree Nearest Neighbor Search

Complete the code below. For each /* n */ , write the missing recursive call.

```
// in-order: axis = depth % 2
void nearest(kdnode *root, double target[2],
            int depth, kdnode **best, double *bestDist) {
    if (root == NULL) return;
    int axis = depth % 2;
    double d = (root->point[0]-target[0])*(root->point[0]-target[0])
        + (root->point[1]-target[1])*(root->point[1]-target[1]);
    if (d < *bestDist) {
        *bestDist = d;
        *best = root;
    }
    if (target[axis] < root->point[axis]) {
        /* 1 */ nearest(root->left, target, depth+1, best, bestDist);
    } else {
        /* 2 */ nearest(root->right, target, depth+1, best, bestDist);
    }
    double diff = target[axis] - root->point[axis];
    if (diff*diff < *bestDist) {
        /* 3 */ nearest(
            target[axis] < root->point[axis]
                ? root->right : root->left,
            target, depth+1, best, bestDist
        );
    }
}
```

#1 nearest(root->left, target, depth+1, best, bestDist);

Explanation (4.1): Recurse into the subtree on the same side as the target.

#2 nearest(root->right, target, depth+1, best, bestDist);

Explanation (4.2): If the target coordinate is greater, we explore the right subtree first.

#3

```
nearest(
    target[axis] < root->point[axis]
        ? root->right : root->left,
    target, depth+1, best, bestDist
)
```

);

Explanation (4.3): If the hypersphere around the current best crosses the splitting plane, we must explore the opposite branch as well.

Problem 5 (6 points): Heap Runtimes

Implementations:

- A: Binary heap
- B: Fibonacci heap
- C: Unsorted array

For each operation, write the letter (once each):

1. $\Theta(1)$ *amortized* insert: \boxed{C}

Explanation (5.1): Appending to an unsorted array takes $O(1)$.

2. $\Theta(1)$ *worst-case* find-min: \boxed{B}

Explanation (5.2): Fibonacci heaps maintain a pointer to the minimum, so it's $O(1)$.

3. $\Theta(\log n)$ *worst-case* extract-min: \boxed{A}

Explanation (5.3): Binary heap extract-min percolates down in $O(\log n)$.

Problem 6 (8 points): Depth-First Search Code Completion

```
void dfs(Graph *g, int v) {
    /* 1 */ visited[v] = true;           // mark v
    printf("%d ", v);
    for (int u = 0; u < g->n; u++) {
        if (g->adj[v][u] && !visited[u]) { // edge exists & not visited
            /* 3 */ dfs(g, u);
        }
    }
}
```

#1 `visited[v] = true;`

Explanation (6.1): We mark the current node as visited to avoid cycles.

#2 In the if clause: `g->adj[v][u] && !visited[u]`

Explanation (6.2): We only recurse if there's an edge and the neighbor is unvisited.

#3 `dfs(g, u);`

Explanation (6.3): Recurse on the neighbor u.

Problem 7 (15 points): AVL Tree Traversals

Insert and then give preorder and inorder of the resulting balanced tree.

- a) Insert {30, 20, 10}:

Preorder: 20, 10, 30

Inorder: 10, 20, 30

Explanation (7.a): A right–right imbalance at 30–20–10 triggers a single right rotation at 20.

- b) Insert {10, 20, 30}:

Preorder: 20, 10, 30

Inorder: 10, 20, 30

Explanation (7.b): A left–left imbalance at 10–20–30 triggers a single left rotation at 20.

- c) Insert {20, 10, 30, 25, 27}:

Preorder: 20, 10, 27, 25, 30

Inorder: 10, 20, 25, 27, 30

Explanation (7.c): Insertion of 27 under the right subtree of 20 creates a left–right case at 30; fix with right rotation at 30’s left child (25), then left rotation at 20.

Problem 8 (16 points): Dijkstra's Algorithm Code Completion

```
void dijkstra(Graph *g, int src) {
    int dist[MAXV];
    bool known[MAXV] = {false};
    /* 1 */ for (int v = 0; v < g->n; v++) dist[v] = INT_MAX;
    dist[src] = 0;
    for (int i = 0; i < g->n; i++) {
        /* 2 */ int v = minVertex(dist, known, g->n); // find unknown min
        known[v] = true;
        for (Edge *e = g->adj[v]; e != NULL; e = e->next) {
            int w = e->to;
            int wgt = e->weight;
            /* 3 */ if (!known[w] && dist[v] + wgt < dist[w]) {
                /* 4 */ dist[w] = dist[v] + wgt;
            }
        }
    }
}
```

#1 for(int v=0; v<g->n; v++) dist[v]=INT_MAX;

Explanation (8.1): Initialize all distances to “infinite.”

#2 int v = minVertex(dist, known, g->n);

Explanation (8.2): Select the nearest unknown vertex.

#3 !known[w] && dist[v] + wgt < dist[w]

Explanation (8.3): Check if going through v gives a shorter path to w.

#4 dist[w] = dist[v] + wgt;

Explanation (8.4): Relax the edge by updating w's distance.